

Discovery-Question-Bank

 About this document ▾

This document provides a list of questions generated with Claude that can be used for conducting discovery interviews with client stakeholders.

Discovery Interview Question Bank

A resource for conducting deep dive, stakeholder interviews in the discovery of how to approach artificial intelligence training. The training is not just about teaching how to use AI as a skill, but also how to develop literacy an fluency with AI. The **critical piece** is learners developing the ability to identification where and how to use AI, including AI Agents. I am not teaching you how to build AI Agents, I am teaching you to identify where to use AI Agents.

Using this question bank

This bank is designed to follow the arc of an interview from open to close. With reflection of who your client is, review the question bank days in advance and select questions you may ask.

Each question has next to it a tag for who it is meant to be asked to, as well as, some insight into what you should listen for and response.

Tag	Who
All	Any stakeholder
Leader	Sponsor, executive, people manager
Learner/Exemplar	Incoming learner, exemplar employee
Tech	IT, security, governance, or platform owner

Additionally, some questions have specific markers to identify what they are meant to hone in on. These markers are aligned with learning more about the client and AI.

Mark	Meaning
◆	Agent opportunity. Most likely to surface where an agent could actually live.
▲	AI literacy and fluency. Reveals where this person or their team actually sits, through evidence rather than self-rating.

Use a table below to skip ahead two relevant sections for specific stakeholders.

Stakeholder	Lean into
Sponsor / executive	§2, §9, §11, §12, §13
People leader	§3, §4, §5, §8, §9, §11
Learner / exemplar	§3, §4, §5, §6, §11, §12
Technical / governance	§7, §10 — and little else; keep it to 25 minutes

§1 • Opening

Open the interview with a statement. This will set the tone. For example:

Thanks for being here today and we're really happy to meet with you about this training that we're preparing to design. What we really hope to get today is a broad sense of what your work looks like, what you're struggling with, then how we can make your life better working with AI. So I have a series of questions I'd like to go through and please be open and honest as you can. There's nothing in here that will go back to your leadership. It is simply to help us make a strong plan for the training. That said do you mind if we record this call?

Your first question should be warm-up, so pick just one from below.

Question	For	Listen for
How long have you been at [company] , and what do you do in your role?	All	Tenure shapes whether they see the process as fixed or changeable
What's a normal day look like for you?	All	Rhythm, interruption load, meeting density

Question	For	Listen for
What did you hear about this training before we scheduled the call?	All	The internal story being told about the program — often not the one Sales told you

§2 • Why This, Why Now

The following questions are designed to be asked of leaders. It is best to explore one or two of these early on as it will frame everything that comes after. Do not use these questions with individual contributors or technology people.

Question	For	Listen for
<u>What business pressure is sitting behind this training request?</u>	Leader	Revenue, margin, headcount freeze, a client escalation, a board slide
<i>Follow up:</i> Is there a number attached to the motivation?	Leader	Whether a target exists or this is directional
Did something happen that made this training a good idea to a funded one?	Leader	The originating incident. Usually specific and usually revealing.
Why at this particular moment has this training become important to you?	Leader	Deadline pressure, competitive move, budget cycle
Who else had to say yes for this to happen?	Leader	The hidden stakeholders you haven't interviewed
<u>What have you tried before on this subject and what happened? OR Have you tried to do this internally and what happened?</u>	Leader	Don't repeat a failed intervention
<i>Follow up:</i> What did people say about it afterward?	Leader	Institutional memory that will walk into your room

§3 • Role, Team, and Orbit

These questions are designed to help understand the shape of the team, including who is in it and who are at the edges. They are interested in how work gets passed along and what handoffs look like.

Question	For	Listen for
<u>What does your team own end-to-end?</u>	All	Scope boundaries and accountability
<u>How many people are in your team and what are their roles?</u>	All	Whether "the team" is 2 people or 40
◆ Who hands work to you, and who do you hand work to?	All	Handoff seams — where agents live
<i>Follow up:</i> What does this hand-off look like? Does anything interfere?	All	Understand further what an agent may do
Who outside your team do you depend on to get your work done?	All	Cross-functional friction, approval bottlenecks
◆ <i>Follow up:</i> How long do you typically wait on them?	All	Queue time for understanding an automation target
◆ <i>Follow up:</i> <u>Do you know what their bottlenecks may be and are they similar to yours?</u>	All	Common stress points where a broad solution may live.
◆ <u>What is the weekly rhythm like from Monday through Friday?</u>	All	Scheduled work is agent-eligible work
Where do you and your team actually talk to each other — Slack, Teams, email, meetings?	All	Where an agent's output would need to land to be seen

§4 • The State of Work

Anchor questions in **a recent event** to get a specific direction. If pulling from general memory, then it is likely to be inaccurate and generalized.

Question	For	Listen for
I'd like you to think back over your last full week. What took most of your time?	All	An instance you can walk through
◆ <u>Let's take that one big drag. Will you walk me through it from beginning to end?</u>	All	Step count, systems crossed, data formats. The foundational agent question.
<i>Follow up:</i> What kicks it off?	All	Scheduled, event-driven, or someone asking
<i>Follow up:</i> Who else touches it along the way?	All	Handoffs and approval gates

Question	For	Listen for
<i>Follow up:</i> Where does it stall or come back to you?	All	Rework loops and failure modes
<i>Follow up:</i> While you are doing it, what apps do you have open on your screen?	All	The real stack, not the stated one
◆ <i>Follow up:</i> How long does one full pass take?	All	Pair with frequency to size the opportunity yourself
Walk me through the actual clicks of [specific core task].	Learner/Exemplar	Task analysis at real grain size
◆ What part of that work is judgment, and what part is mechanical?	All	The automation line, in their own words
◆ Which part would you never want a machine touching? Why?	All	Human-in-the-loop boundaries. One of the highest-yield questions in the bank.
◆ What does "done well" look like here? Who decides that or how is it decided?	All	Quality criteria — an agent needs a success condition and this is it
◆ When this workflow goes wrong, how does it go wrong?	All	Failure modes determine how much autonomy a design can grant

§5 · Agent Opportunity Signals

Design note: Don't ask stakeholders where agents could be used. Your client's ability to identify where they could be used *is the skill you're teaching*. And if they could already do it, they wouldn't need the training. Ask about the **characteristics of the work** and do the identification yourself afterward.

Question	For	Listen for
◆ What do you copy and paste between systems?	All	Copy-paste is a workflow confessing where its seams are
◆ <i>Follow up:</i> At what point does something get exported to a spreadsheet?	All	Export = the system boundary an agent would bridge
◆ Where does the data live, and where does the answer need to end up?	All	The gap between those two is the agent

Question	For	Listen for
◆ What runs on a schedule — every Monday, month-end, quarterly?	All	Frequency justifies build effort
◆ <u>What gets done the same way every time, where you could write the rules down?</u>	All	Rule-following work is the highest-yield first agent; it is also a potential Skill
<u>Is there work where two people would produce different results from the same inputs?</u>	All	Divergence = judgment (keep human) or inconsistency (standardize first, then automate)
◆ <u>Where does the same work get done differently by different people on your team?</u>	Leader	Standardization opportunity. Leaders see this; Individual Contributors can't.
◆ What's the work you'd never hire someone full-time to do, but wish got done?	Leader	Latent demand — often the best capstone project
◆ What do you check or monitor that rarely turns out to need action?	All	Classic monitoring-agent candidate
◆ What question do people ask you over and over?	All	Retrieval-agent candidate; also a job-aid candidate
<i>Follow up:</i> Where would someone find that answer themselves if they knew where to look?	All	Whether the knowledge is bounded and findable
◆ If a draft of that came back wrong, who'd catch it and how?	All	Reversibility. Sets the autonomy ceiling.
What use cases do you already anticipate for this?	All	Their framing. A starting point, not a finding.

§6 • Current AI Reality

Identify where this person and their team actually are with AI today. Measure through **instances and artifacts**, and never self-rating.

Question	For	Listen for
What AI tools do you have access to today?	All	Stated access — verify against §7 and §10
▲ Tell me about the last three times you actually used one.	All	This is the real measurement
▲ <i>Follow up:</i> Did you use any tools in combination, and if	All	Knowledge of tools and knowledge of how to pick them

Question	For	Listen for
so, how did you decide that?		for tasks
▲ <i>Follow up</i> : What were you trying to get done?	All	Task-level vs. workflow-level thinking
◆▲ <i>Follow up</i> : What did you do with the output?	All	Whether AI output feeds a downstream step or dead-ends. The downstream step is the agent.
▲ <u>What's something you tried that didn't work?</u>	All	Failure stories reveal the ceiling faster than success stories
▲ Has anyone on your team <i>built</i> something — a project, a Skill, a custom GPT, an automation?	All	Separates chat users from builders
▲ <i>Follow up</i> : What did they build it for and did it have the intended impact?	All	Designing AI solutions for impact
▲ <i>Follow up</i> : Can you show me, or send it over?	All	Artifacts beat descriptions every time
▲ How do people use the word "agent" used around here?	All	Shared vocabulary, or the lack of it. Expect wide variance.
▲ What is your team's attitude towards AI right now?	Leader	Understand the lay of the land and their frustration
<i>Follow up</i> : If they are showing some resistance to adopting AI, what is that driven by?	Leader	Gather reasons they hear or believe their teams are lagging
▲ Who really stands out on your team with this? Why?	Leader, Learner/Exemplar	Names your near-peer exemplar
▲ Has anyone used, or are they using, AI in a way that surprised you?	Leader	Shadow adoption, and sometimes shadow risk
▲ How are you seeing your team use AI in their work right now?	Leader	The leader's-eye view. Compare against what participants told you — the gap is often large and always informative.
<i>Follow up</i> : Is there anyone you'd like to see using it <i>more</i> ?	Leader	Under-adoption is as common as over-adoption, and leaders rarely volunteer it

§7 • Tools and Tech Stack

Establishes what is actually available, actually connected, and actually permitted from IT. Understand the state of IT understanding with the real work happening.

Question	For	Listen for
<u>What sort of access do teams at your company have to AI tools?</u>	All, Tech	What they are able to use and how it matches with their work.
<i>Follow up:</i> Are some blocked or on hold? Why is that?	All, Tech	ITs rational for blocking and if there is frustration from others
Exact license tier and enabled features, per participant group?	Tech	"They have Claude" and "they have Claude with X enabled" are different programs
<i>Follow up:</i> What's disabled, and what would it take to turn it on?	Tech	Lead time. Audit logs and scoped folder access are common blockers.
▲ <u>When you have a choice between [Tool A] and [Tool B], how do you decide?</u>	All	Whether a mental model exists or it's habit — tells you if tool-selection needs teaching
<u>Which tool is the sanctioned path for which kind of work?</u>	Tech, Leader	Tool-blending can only be taught inside sanctioned combinations
◆ What integrations or connectors are live between systems today?	Tech	Whether orchestration labs are real or hypothetical
◆ What's the system of record for [the data in question]?	Tech, All	Where an agent would have to read from
<u>Can our instructor get an account and demonstration in your environment before we build?</u>	Tech	Non-negotiable ask. Design blind and you will rebuild.
<u>Is there de-identified or synthetic sample data we can use?</u>	Tech	Lab realism without the compliance fight
Network, VM, or whitelisting needed on delivery day?	Tech	Delivery-day failure prevention
What's on the roadmap to change in the next two quarters?	Tech, Leader	Don't teach to a stack that's being replaced
<u>Are there AI tools that you are learning on your own outside of work? What draws you to them?</u>	All	Gauge overall interest and willingness to explore

§8 • Pain, Blockers, and Overwhelm

Find out what part of the work currently hurts, is considered inefficient, or just frustrating. Uncover anything standing in the middle and acting as a blocker.

Question	For	Listen for
◆▲ If you could hand one recurring thing to a capable assistant tomorrow, what is it?	All	Their own framing of the opportunity
<i>Follow up:</i> What would you do with the time gained?	All	Whether the payoff is capacity, quality, or strategy
▲ What's blocking you from using AI more than you do now?	All	Sort the answer: skill, access, permission, trust, or time . Only one is a training problem.
Where do you personally feel most underwater right now?	Leader	Personal framing gets a truer answer than "where is the team overwhelmed"
What have you watched happen on your team that made you think <i>we need to fix this</i> ?	Leader	Concrete, observed, and usually vivid
What's the thing your team complains about that you can't currently solve?	Leader	Where training could deliver a visible win
◆ What's something people asked for that IT or engineering said no to?	All	Blocked demand — often the best lab material
◆ What are you doing manually that you assume must have a better way?	All	Awareness of the gap without knowledge of the solution — exactly your audience
What's your biggest blocker in this work generally?	All	Deliberately broad. Sometimes the answer isn't about AI at all, and that matters.
Do you have any people acting as blockers to AI tools? What is motivating that?	All	Uncover resistance you may encounter and rational

§9 • Capability Strategy

Two halves of the same question. §9a asks how far up they want people to climb. §9b asks what has to stay with the human on the way up. Leaders almost always hold both at once, and the second one they raise more reluctantly.

§9a • The Skill Ladder

How much thought has gone into AI capability as a **people strategy** rather than a tool rollout. Asking this well often produces an artifact — and a client-authored capability ladder is the single most useful document you can be handed before design.

Design note: If they hand you a ladder, align your learning objectives to their tier names and use their vocabulary in the deck. The training will read as native rather than imported. If they don't have one, the absence is itself a finding worth raising at HLD Validation — a program without a progression model gets evaluated on how the room felt.

Question	For	Listen for
▲ <u>Do you have a defined progression for AI capability — levels you want people to move through? For example, something like: Beginner → Practitioner → Advanced → Agent Builder → Engineer.</u>	Leader	Whether one exists at all. The answer tells you how mature their thinking is, and the example gives them the shape of a good answer without putting words in their mouth.
<i>Follow up:</i> What are the levels called, and what does someone at each one actually do?	Leader	Behavioral definitions vs. label-only tiers. Behavioral ones are usable as objectives.
<i>Follow up:</i> How far would you say you've gotten developing people with this?	Leader	Have they been pursuing this strategy actively already
<i>Follow up:</i> <u>Would you share it with us?</u>	Leader	Always ask. A client ladder becomes your competency map scaffold and your assessment blueprint.
▲ <i>Follow up:</i> <u>Where are most of your people today, and where do you need them by [date] ?</u>	Leader	The gap you're being hired to close, stated numerically
<i>Follow up:</i> Which level does this training need to move people to?	Leader	Direct scope check. Watch for a ladder that tops out below what they bought.
If you don't have one — would a draft be useful to you?	Leader	Turns a gap into a value-add, and gives you license to define the progression yourself
Does AI capability show up anywhere in job descriptions, reviews, or promotion criteria?	Leader	The real test of whether the ladder is strategy or a slide. If it's in performance criteria, adoption is nearly solved. If not, expect motivation to be your problem.

Question	For	Listen for
Who owns AI capability development in your company — L&D, IT, a center of excellence, or nobody?	Leader, Tech	Your ongoing counterpart, and who inherits sustainment
Are there AI champions or a community already? What are they doing?	Leader	Existing infrastructure you can plug the program into rather than duplicate

§9b • Quality, Judgment, and Offloading

Leaders commonly want more AI capability on their teams *and* privately worry that more use means less thinking. Both are true at once and they rarely raise the second one unprompted — it sounds like distrust of their own people. Ask, and ask without judgment, because the answer defines the quality bar the training has to teach to.

Design note — the tension is yours to resolve, and saying so builds credibility. If a leader is worried about cognitive offloading and you sell agent-building, you may be teaching the thing they fear unless the judgment layer is explicit in the design: when to verify, how to critique output, where the human decision sits, what "I checked this" actually means.

Question	For	Listen for
How are you feeling when you see someone produce some excellent work partnering with AI?	Leader	Genuine excitement or anxiety
▲ Are you satisfied with the quality of work being done with AI?	Leader	Opens the door. "Mostly" and "it's fine, but..." are both invitations — take them.
<i>Follow up:</i> Tell me about a time it could be better.	Leader	Convert the anxiety into an incident. A real story of AI-assisted work landing badly is case-study material.
▲ When someone brings you AI-assisted work, can you tell where their thinking is in it?	Leader	Concrete and observable, unlike asking about "agency." The answer is usually a rueful no.
Are you concerned about people offloading their expertise to AI rather than applying it alongside?	Leader	The core worry, stated plainly. Naming it first gives them permission to agree.
<i>Follow up:</i> Have you addressed this with your teams?	Leader	Willingness to engage with their teams about AI
<i>Follow up:</i> Does that concern differ for people early in their career vs. experienced people?	Leader	Sharpest for juniors, who haven't built the expertise to offload yet. Directly shapes design for early-career cohorts.

Question	For	Listen for
◆ <u>What are the parts of your team's work where the thinking has to stay with the person?</u>	Leader	Leader-level twin of the §4 human-in-the-loop question. Scopes what agents may and may not own.
▲ <u>How do people check AI output before it goes out the door? Is there a standard or workflow?</u>	Leader, All	Verification behavior is the observable proxy for judgment. If there's no norm, that's a module.
Is there an expectation that people flag when something was AI-assisted?	Leader	Disclosure culture. Tells you how openly people can work in the room.
<u>Is there a version of AI adoption on your team that would worry you?</u>	Leader	Invites the fear directly without asking them to accuse anyone
<u>What would you want people to be better at, not just faster at?</u>	Leader	Reframes from efficiency to capability. Often the most useful sentence in the interview.

§10 • Strategic Direction and Ceiling

Mostly leadership territory. This is where you find out whether the training's ambition matches what the organization will permit.

Question	For	Listen for
<u>Picture your team six months after this training. What should they be doing that they aren't doing now?</u>	Leader	Behavioral outcome — raw material for learning objectives
What would this unlock for the team that isn't possible today?	Leader	Strategic framing beyond time savings
<u>What are people allowed to build on their own, and where does engineering have to be involved?</u>	Leader, Tech	The permission ceiling. Draw this line before designing labs.
<i>Follow up:</i> Who decides which side of that line something falls on?	Leader	The escalation path participants will need to know
<u>If someone builds something useful in this training, what happens to it afterward?</u>	Leader	Whether there's a path to production or it dies on the laptop
<u>What would you <i>not</i> want your team automating, even if they could?</u>	Leader	Values and risk appetite. Often about client relationships.
If we came back and said the plan should shift — more of X, less of Y —	Leader	GA requires this scope-check in Discovery. Surfaces flexibility

Question	For	Listen for
how would that land?		before HLD Validation.

§11 • Permission and Governance

Determines what you are allowed to teach. Explore these before you commit to a lab idea because a governance surprise late in development costs a rebuild.

Question	For	Listen for
What is not allowed to go into an AI tool? Client data, PII, contracts, source code?	All	Lab data design
<u>Is there a written AI policy? Who owns it?</u>	Leader, Tech	Absence of one is a red flag, not a green light
<i>Follow up:</i> Is it a living document? How often are you updating it?	Leader, Tech	Maintenance of AI governance policy
<i>Follow up:</i> Can we see it and use it in the training?	Tech	Read it before you design. Assume nothing.
<u>How do people currently find out what's allowed?</u>	All	If the answer is "they ask around," governance is a training topic too
Has anything been switched off that people wish were on?	All	Feature gaps you'll hit mid-lab
<u>Is anything under review right now that might change before delivery?</u>	Tech	Moving-target risk
<u>What happens if someone uses a tool they shouldn't?</u>	Tech, Leader	Severity of consequence shapes how boldly you can teach

§12 • The Training Experience

People, Safety, and Energy

Intelligence about the cohort itself. Who is walking into the room, what they can say out loud, and what will make them disengage. Learners answer these best.

Question	For	Listen for
▲ <u>Across the people attending, how wide is the gap between most and least capable?</u>	Leader	Audience heterogeneity — the biggest design risk in mixed cohorts. Pairs directly with §9.

Question	For	Listen for
<i>Follow up:</i> Do those groups normally work together?	Leader	Whether mixing them creates collaboration or silence
Who's going to be skeptical in the room, and why?	Leader	Emotional arc and resistance planning
Is attendance voluntary or assigned?	All	Changes the opening module completely
▲ Who in your group is the go-to person for this? Why them?	Learner/Exemplar	Get the name. The instructor can call on them deliberately.
▲ <i>Follow up:</i> What's something a colleague does that you wish you knew how to do?	Learner/Exemplar	Peer-learning targets for a share-out
How comfortable is this group sharing mistakes out loud?	Learner/Exemplar, Leader	Psychological safety read. Determines public debrief vs. paired work.
<i>Follow up:</i> Will their manager be in the room?	Learner/Exemplar	Changes disclosure entirely. Design accordingly.
What's your red flag in a live session — the thing that makes you open your email?	Learner/Exemplar	Anti-patterns to design out. Usually lecture length or demos with no hands on keys.
<i>Follow up:</i> Best live training you've ever attended — what made it work?	Learner/Exemplar	Their working definition of good. Cheap to ask, occasionally gold.
What's a real horror story with [topic] we should solve in class?	Learner/Exemplar	Scenario material with a built-in emotional hook
If you had one expert for ten minutes, what would you ask?	Learner/Exemplar	Seeds the Q&A; sometimes reveals a missing module
Why are you attending?	Learner/Exemplar	Volunteered, volun-told, or curious
<i>Follow up:</i> What do you think this is going to be?	Learner/Exemplar	Expectation gap to close in the opening

§13 • Success, Measurement, and What Comes Next

How they will judge the program, whether anyone is positioned to measure it, and what happens to the momentum once the training is over.

Question	For	Listen for
<u>How would you know in 90 days whether this worked?</u>	Leader	Metrics — and whether any exist
<i>Follow up:</i> Is anyone tracking that today?	Leader	Baseline availability. Usually no.
What would make you say this was worth [N] days of your team's time?	Leader	Success definition in their words
◆ What's one specific thing on your plate next week that this could help you finish?	All	The single best source of authentic lab material. Ask nearly everyone.
<u>What do you feel that need to walk away with — a checklist, a working agent, a peer contact?</u>	Learner/Exemplar	Take-away format. Shapes the participant guide.
Once the training ends, who owns keeping the momentum going?	Leader	Sustainment. Usually nobody, and that's worth naming out loud.
<u>Is there a place where people would share what they built?</u>	All	Channel, community of practice, or a gap you should recommend filling

§14 • Close-out

The last five minutes. Use this to secure the artifacts you need.

Question	For	Listen for
<u>What would make this training a waste of time?</u>	All	The failure mode to design against. People answer this more honestly than the positive version.
<u>What haven't I asked that I should have?</u>	All	Costs nothing, occasionally reframes the whole engagement
Who else should we be talking to?	All	Snowball toward the exemplar employee
<u>Can you send us</u> [a sample deliverable / a redacted file / a screenshot of the system / your capability ladder] ?	All	Artifacts. Ask on every call.
Is it alright if we follow up with a question or two by email?	All	Cheap option on a second pass

Appendix A — The Probe Ladder: task → workflow → agentic

Stakeholders will often answer at **task** level by default. For example, "I use AI to write emails." This however will not support an agent lab. So ask questions around the activity.

- 1. **Upstream** — "What happens before that email gets written?"
- 2. **Downstream** — "And after you send it, what happens next?"
- 3. **Handoffs** — "Who else is in that chain?"
- 4. **Cadence** — "How often does the whole thing run, start to finish?"
- 5. **Judgment** — "What decision gets made in there, and on what basis?"
- 6. **Failure** — "When it goes wrong, how does it go wrong?"

Stop when you have all six: trigger, steps, systems, handoffs, decision points, failure modes. That's a workflow. Anything less is a task.

Appendix B — Low-Yield Questions

Questions that may feel productive, but really are not. For each one, try to convert is into a more meaningful question. Some examples:

Low-Yield	Because	High-Yield
"Rate your AI proficiency 1–5"	Uncorrelated with capability.	"Tell me about your ability to get what you needed the last few times you used an AI tool?"
"What do you already know about [topic]?"	Self-assessment against a standard they don't have. Novices will overestimate, and experts will hedge.	"Tell me about the last three times you used one."
"What do you hope to learn?"	Produces a wish list bounded by what they can only imagine. If they know, then they could learn it themselves.	"What's is commonly on your plate that learning this could help with?"
"Do you think AI could help with that?"	Leading. They will say yes, but often they have no idea how. They are looking to say "yes".	"What part of that is judgment vs. mechanical?"
"Where could you use an AI Agent?"	Same problem but worse. Part of the problem is that the client doesn't know where they can do that.	Any ♦ question in §5

Low-Yield	Because	High-Yield
"How much time would that save?"	Made-up numbers that are just estimations from one person to the next.	Ask frequency and duration separately; do the math yourself
"What works better with your people, individual work or breakout groups?"	Asks learners to make a design decision they have no basis for; answers track personality, not learning	"How comfortable is this group sharing mistakes out loud?"
Any two-part question	They answer the second half only.	Divide it into two questions or drop one part all together.
Anything answerable with "yes"	Closed questions are a dead end.	Open with <i>walk me through</i> or <i>tell me about</i>

Appendix C — Post-Interview Capture

Within 24 hours after the interview capture some or all of the following into notes:

- **Role and team shape** — size, roles, reporting line
- **Workflows** — trigger → done, one line each
- **AI Agent candidates** — sorted by who could build them
- **Capability ladder** — document theirs, if they have one
- **Quality bar and offloading concerns** — what "good" means to the leader
- **Offloading concerns** — what they do not want automated away
- **Tools and access** — stated versus verified
- **Constraints and red flags** — governance, data, disabled features
- **Verbatim quotes** — the ones stand-out
- **Named people** — the 'go-to experts'
- **Artifacts requested** — who is the owner and when to follow-up
- **What this changes** — capture what we heard and what we recommend
- **Questions for the next interview** — the thread you still need to pull on